

Compilers Final Project - Concolic Testing in LLVM

Carl Hildebrandt
The University of Virginia
ch6wd@virginia.edu

Will Leeson
The University of Virginia
wel2vw@virginia.edu

ABSTRACT

Concolic Execution is a testing strategy that utilizes both concrete testing and symbolic execution to make up for the shortcomings of each technique. Concrete execution is as fast as the program in question. However, if inputs are being generated at random, it may take far too many tests to satisfy all paths of execution. Symbolic execution can find all paths of execution. However, this can be computationally expensive. Concolic execution leverages one against the other. Symbolic execution can find the paths in which concrete testing has difficulty reaching. Concrete execution can execute the hard to solve paths. Many concolic execution tools target a specific language. This can lead to time spent re-targeting tools to other languages. LLVM is a compiler infrastructure that allows for creation of compilers that produce the same intermediate representation, LLVM bytecode. In this paper, we design a concolic execution strategy that can be used on LLVM bytecode. This will allow for a concolic execution tool that can be used on languages that can be compiled into LLVM bytecode.

KEYWORDS

Concolic Execution, LLVM, Static Analysis, Dynamic Analysis, Compilers

ACM Reference Format:

Carl Hildebrandt and Will Leeson. 2018. Compilers Final Project - Concolic Testing in LLVM. In *UVA '19: CS6161 - Compilers, December 06, 2019, Charlottesville, VA*. ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/1122445.1122456>

1 INTRODUCTION

Symbolic execution and concrete execution are two testing techniques that can be used to find and test execution paths through a program. Symbolic execution finds the constraints of a program to discover paths of through the program. In order to test these paths, it generates the constraints that need to be satisfied to execute it. It then passes these constraints to a tool, like the SMT solver Z3 [3], which decides if these constraints are satisfiable. If it is satisfiable, it can model this solution with an input that will execute the path in question. Solving these constraints can quickly become difficult as the constraints become more and more complex. As a result, using symbolic execution tools can become infeasible with complex programs.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
CS6161, December 06, 2019, Charlottesville, VA

© 2018 Association for Computing Machinery.
ACM ISBN 978-1-4503-9999-9/18/06...\$0.00
<https://doi.org/10.1145/1122445.1122456>

On the other end of the spectrum, there is concrete execution. Unlike symbolic execution, concrete execution is relatively inexpensive. The test is as expensive as the program in question. This allows for running many tests to exercise as many paths as possible. Often the input for automated concrete execution is randomly generated. If the probability of going down any path is fairly uniform, testing all paths is fairly simple. However, if the probability of hitting certain paths is low, it can be difficult to exercise all paths. However if combined correctly, these tools can compliment each other resulting in a more accurate yet cheap automated test tool.

Another issue common in many software testing tools is that the tool can only target a specific language. This can lead to several tools that may accomplish the same task, only for different targets. The scope of the tool can be expanded if it targets a more general language. LLVM is a compiler infrastructure that provides the tools for developers to make compilers for their own languages or existing languages. LLVM compilers produce an intermediate representation known as LLVM bytecode. Tools that target LLVM bytecode should then be able to operate on the various languages that have a compiler that produces LLVM bytecode.

The LLVM infrastructure allows for various "passes" to be run on LLVM bytecode. These passes can perform static analysis, such as dependency analysis, which finds dependencies in memory usage. They can also transform code by adding instructions or optimizing it with passes like constant propagation [6]. We create two passes, one that statically collects information about the program at compile time, and one that instruments code to output information dynamically at execution time. Using the information generated by these passes, we can generate input that will explore the paths of a program more efficiently than symbolic execution and more exhaustively than concrete execution.

We present a framework to do concolic testing on LLVM bytecode and show the workflow of using this framework to exercise the paths of execution of a program. This framework uses the two LLVM passes written by the authors. The first pass is a static pass run at to generate the constraints of a program. The constraints are used to ensure we exercise all paths. The second pass is a dynamic pass that instruments the code that instruments the bytecode to report the trace of the program with respect to the constraints it reaches and whether or not the input satisfies them. Using these traces, we can generate new input that follows paths not executed by past input. More specifically this paper has the following contributions:

- (1) We design a concolic test engine targeted at LLVM bytecode.
- (2) We design and implement a dynamic LLVM pass that instruments a program to report a trace of that programs execution.
- (3) We design and implement a static LLVM pass that analyzes and displays all constraints in a program.
- (4) We design and implement a python program that uses an SMT solver to solve a given set of constraints.

2 MOTIVATION

As software systems get larger, symbolic execution is less likely to be able to perform an efficient, correct analysis to generate inputs that tests all paths of execution. At the same time, concrete execution using a fuzzing tool that generates random input will not be able to exhaustively test these systems. There is no guarantee that a fuzzer can exercise all paths in a program without generating all possible input.

If a fuzzer were to analyze Code Snippet 1, it could generate input that could easily exercise the path $(input > 5) \wedge (input > 6)$ with any input greater than 6. To execute the path $(input > 5) \wedge \neg(input > 6)$, there is only one input, namely 6, that would satisfy this constraint. The randomness of the fuzzer doesn't guarantee this path will be found until all inputs have been tested. If this input is an int and this were a C program, this in the worst case situation could take 2^{32} executions. However, a symbolic execution tool could solve the constraints of the path, and simply return three tests that would satisfy all paths of the code snippet, for example: $model((input > 5) \wedge (input > 6)) = 7$, $model((input > 5) \wedge \neg(input > 6)) = 6$, and $model(\neg(input > 5)) = 0$.

Code Snippet 1: Possible Error Code

```

1  ...
2  if (input > 5) {
3      ...
4      if (input > 6) {
5          ...
6      }
7      else {
8          error();
9      }
10 }
11 ...

```

A problem arises with scalability when using symbolic execution tools. If a symbolic execution tool were to analyze a large codebase, the constraints to execute path with compound and may reach a point where solving the constraints may be more computationally expensive than fuzzing the program. If the paths that are computationally complex to solve are executed with a reasonable high probability, then a fuzzer can be used to generate random input that will hit this path instead.

Using these two types of tools together can lead to avoiding both situations described above. In this paper we design a workflow that would allow for concolic testing of LLVM bitcode. Concolic testing uses concrete execution combined with symbolic constraints to generate new inputs which explore a program.

3 BACKGROUND

Concolic execution is a useful technique that uses both the concepts of concrete execution and symbolic execution. Concrete execution is a fast testing strategy that relies on the input it is given to be effective. Generally, it can take random input to hopefully cover all paths in question. However, this can lead to issues if the input distribution doesn't cover all paths uniformly. If there are several paths a program could take and one of them is taken on only one input out of the entire input space, this can lead to an issue where that path is never taken and is therefore never tested. Symbolic execution on the other hand calculates all paths a program can take. This prevents this issue of missing a path. Unfortunately, symbolic

execution can be very expensive. As path conditions become more complex, it is harder to solve the constraints that would produce the input to follow this path. With concolic execution, the advantages of these two types of tools can be leveraged against each other to produce a faster tool that checks more program paths.

Recent developments in concolic testing have looked into optimizing the process of switching between the concrete and symbolic sides of the tool [11]. By using the probability of a path being taken along with the cost of solving the constraint of paths, a tool can find the optimal time to evaluate a path using the concrete execution or the symbolic execution. Deciding when to choose between these tools is still an open problem. Currently, many tools, like KLEE [2], use various heuristics to decide when to proceed symbolically or concretely.

In recent years, tools like Badger [9] and Driller [10] have combined symbolic execution with fuzzing tools. Fuzzers generate random input that should exercise the paths of execution of a program. Since these tools may not exercise paths with a low probability of execution, combining them with symbolic execution tools creates a concolic testing engine.

4 PROBLEM

The goal of this project was to design a concolic testing engine to test programs which are in an LLVM bitcode format. To accomplish this, we first needed to create a transformation pass in LLVM which transforms the program such that when you run it, symbolic constraints are generated showing the execution of that program run. Using these symbolic constraints a test generator would generate new input which executed a new path in the program. This approach could be repeated until all paths of the program were explored or a timeout value reached.

For our problem the LLVM bitcode was generated by two compilers namely, TIPC and Clang. TIPC is a compiler which takes TIP[1] programs and converts them to LLVM bitcode [8]. Clang is a compiler front end which uses the LLVM compiler infrastructure to generate LLVM bitcode for, among others, C and C++ programs. Once the LLVM bitcode was generated, it would have to have LLVM passes run on it to generate the information we need at both compile and run time. This information will then be given to an SMT solver which will generate new inputs to exercise new paths.

5 APPROACH

The main objective of our approach was to determine the symbolic constraints of a program's execution path. The generated symbolic constraints give us a way to generate input that will exercise the paths of the program. Using this information our technique aims to generate more tests. We generate more tests by looking at the constraints satisfied by a trace and negating the most recent constraint that hasn't been negated to generate test input that will take a different path. This process can be repeated until all satisfiable paths have been explored.

The workflow of our technique is shown in Figure 1. Our approach starts by taking the original program and passing it through our transformation. This transformation is described below in Section 5.1. The transformation instruments the original program with

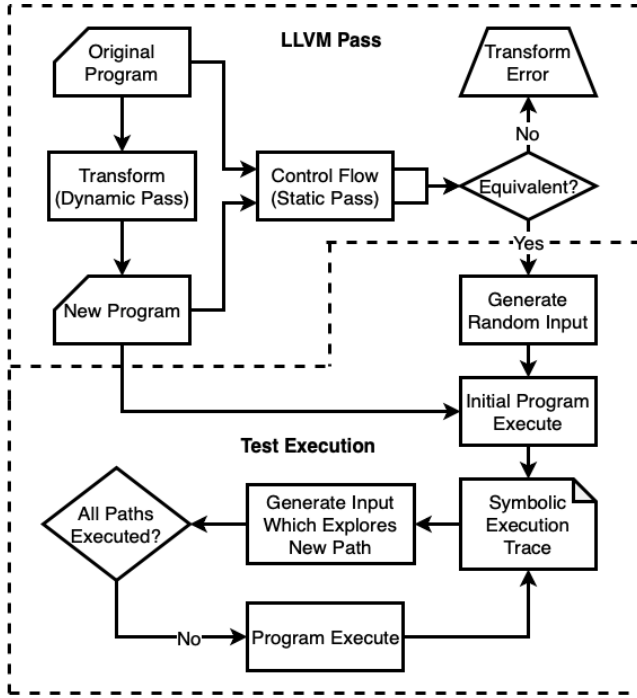


Figure 1: The general overview of our approach. Our approach starts by using LLVM’s pass infrastructure to instrument the original program. We then execute the program using the symbolic execution trace to generate new inputs which drive the program on new paths.

new instructions which evaluate and display the symbolic constraints of a program during execution. In order to determine if the instrumentation resulted in a program which was equivalent in terms of the control flow, a static control flow pass was also instrumented. The static control flow pass is described below in Section 5.2. The static pass went through a program and displayed all flow insensitive symbolic constraints from the program. This was done a check to ensure that dynamic transformation worked as expected.

We then executed the new program using random input. Using the symbolic constraints generated trace from the execution of the program using random inputs, we generate new input for the program. The new input is generated to explore new paths in the program under testing. We repeat this process until all paths have been executed. The final results will be a set of tests which explore different program paths and a set of constraints for the program describing each path through the program.

5.1 Dynamic Instrumentation Pass

Our program transformation works as described in Algorithm 1. The algorithm takes an input P which is the original program we want to run concolic testing on. The algorithm outputs P' : a new program which has been instrumented to output symbolic constraints of each predicate, the current variable values in each predicate and, whether the predicate final evaluation. The goal is

Algorithm 1: Dynamic Pass General Overview

```

Input :  $P$ 
Output :  $P'$ 
1 about_to_branch = False;
2  $P' = P$ ;
3 for Each  $F$  in  $P$  do
4   for Each  $BB$  in  $P$  do
5     for Each  $I$  in  $P$  do
6       if  $I == \text{Load\_Instruction}$  then
7          $LHS = I.get\_LHS()$ ;
8          $RHS = I.get\_RHS()$ ;
9          $match\_vector = (LHS, RHS)$ ;
10      end
11     if  $I == \text{Compare\_Instruction}$  then
12        $op = I.get\_operands()$ ;
13        $op\_m = match\_operands(op, match\_vector)$ ;
14        $P'.log\_before(I, op\_m)$ ;
15       about_to_branch = True;
16     end
17     if  $I == \text{Branch\_Instruction}$  then
18       if about_to_branch then
19         about_to_branch = False;
20          $eval = evaluate\_results(I)$ ;
21          $P'.log\_before(I, eval)$ ;
22       end
23     end
24   end
25 end
26 end
27 return  $P'$ 

```

transform P to P' while ensuring the control flow and decision points of both P and P' are equivalent.

Algorithm 1 is an LLVM transformation pass. It works by first iterating through all functions F , then through all basic blocks BB and then through all instructions I , shown in lines 3, 4 and 5 respectively. We then determine the type of instruction and handle each instruction differently. We are mainly interested in three instruction types namely, the load instruction, a compare instructions and a branch instruction as these instructions are used during program branching. The load instruction is important as it allows us to keep track of how local variables are transferred. The compare instruction compares two local variables, which determine the resulting flow of the program. Thus we are interested in obtaining the symbolic constraints from each of these compare instructions. The final instruction is the branch instruction. These instructions are where the branching actually occur in the program and thus an easy way to determine the result of a compare instruction is to evaluate the input to a branch instruction.

Algorithm 1 checks for a load instruction in line 6. If the instruction is indeed a load instruction the algorithm extracts the left hand side (LHS) in line 7 and right hand side (RHS) in line 8 from the instruction. The local variable in the RHS is being loaded into the

Algorithm 2: Static Pass General Overview

```

Input: P
1 for Each F in P do
2   for Each BB in P do
3     for Each I in P do
4       if I == Load_Instruction then
5         LHS = I.get_LHS();
6         RHS = I.get_RHS();
7         match_vector = (LHS, RHS);
8       end
9       if I == Compare_Instruction then
10        op = I.get_operands();
11        op_m = match_operands(op, match_vector);
12        op_str = op_m.to_str();
13        display(op_str);
14      end
15    end
16  end
17 end

```

local variable in the LHS. Thus we create a vector in line 9 which maps address and variables names of both sides. This mapping is used later during the constraints generation to determine which local variables are being used in each of the comparisons.

The next instruction Algorithm 1 checks for is a compare instruction in line 11. The algorithm starts by getting the operands for the compare instruction in line 12. Using the load instructions vector, each of the operands are mapped to a known variable in line 13. Each of the operands are pass as arguments to a logging function which will display their values during a programs execution. Thus when a program is executed the constraints for each decision point in the program are printed out. Finally, line 15 sets a variable stating that a branch instruction is about to occur to true.

The final instruction which is considered by the algorithm is a branch instruction as seen in line 17. The algorithm is only interested in branches which occur due to decision points such as compares. Other branch instructions are sometimes inserted into a program for instance to call a function. These do not change the flow of the program and so are ignored. Line 19 sets the *about_to_branch* variable to false as the algorithm has found the expected branch instruction. Line 20 evaluates the result of the branch instruction. The result of this branch instruction is also the results of the compare instruction earlier. Finally algorithm 1 prints the result of the branch instruction just before the branch occurs. Thus during execution our new program P' at each decision point, we will print out the constraints containing references to known local variables and the results of each decision point will also be displayed.

5.2 Static Pass

The static pass used an algorithm similar to the dynamic pass and is described in Algorithm 2. There were however two major differences between the dynamic pass and that static pass. The first was that the static pass did not consider branching instructions. That is because the result of branching instructions can only be

Algorithm 3: Input Testing Overview

```

Input, n: P'
1 I = [];
2 for n times do
3   I[n] = random_value();
4 end
5 constraints = [];
6 while I is not NULL do
7   trace = P'(I);
8   constraints.append(trace);
9   I = solveConstraints(trace, constraints);
10 end
11 Function solveConstraints(T, SC)
12   I' = NULL;
13   for Each t in size(T) do
14     new_constraints = T - T[size(T)-t] + !T[size(T)-t];
15     if new_constraints not in SC then
16       I' = smt.solve(new_constraints);
17       break;
18     end
19   end
20   return I';

```

determined at runtime and thus is not applicable to a static pass. The second major difference is that the constraints are converted to a string representation in the static pass and displayed to the user. Thus no evaluation of the constraints takes place when the pass is run.

5.3 Test Execution

The process of generating and executing tests is described in Algorithm 3. To begin executing paths of the program, we first generate n random values, one for each of the augmented program P' 's inputs as seen in lines 2 and 3. Next we create a list which will contain all the symbolic constraints found during program execution in line 5. While we have an input, P' is run with these random inputs in lines 6 and 7. The program execution in line 7 returns a trace containing symbolic constraints for that specific run. This trace is appended to the list of constraints for all runs in line 8. Using the programs current trace and list of previous traces, we call a function solveConstraints which generates an input that will explore a new path in the program. Using this new input, we repeat the process until all paths have been found.

The function *solveConstraints* in line 11 of Algorithm 3 is described as follows. First line 12 sets the input to NULL. If the constraints can not be solved this value will be returned and the concolic testing ended. For each of the constraints in the program trace, we negate that constraint in line 14 by removing the last constraint, negating it and adding it back. If the new constraints have not been seen before, we pass them to an SMT solver which solves them and generates a new input, as shown in line 16.

6 IMPLEMENTATION

Both the dynamic and the static pass were implemented in LLVM 7. In order to test our approaches we built LLVM from source and inserted our pass into the transforms directory of LLVM. Both passes were similar in that they iterated through instructions and identified different instruction types. In this section we will describe how each of the instructions were handled for both passes. We will also describe how the dynamic logging instructions were instrumented into the original program using LLVM. All code is available at: https://github.com/hildebrandt-carl/concolic_testing

6.1 Inserting Logging Instructions

The dynamic pass instrumented a program with logging functions to determine the symbolic constraints evaluated by the program, as well as their value during the program execution. Code Snippet 2 shows a comparison instruction `fcmp` which inherits from the base comparison instruction class. We can see that it checks if the local variable `%2` is equivalent to the floating value `0.0`. The result of that `%cmp` is then used to determine if the program should branch or not.

Code Snippet 2: Original Branching Code

```
1 Inst (fcmp): %cmp = fcmp oeq float %2, 0.0
2 Inst (br): br i1 %cmp, label %if.then, label %if.else
```

Our dynamic pass worked by inserting instructions which made calls to logging functions in a logging library we created. The logging library contained the declaration of multiple logging functions which we could use depending on the circumstances. For instance when a compare instruction was reached we printed out both the symbolic constraint, as well as each of the constraints run time values. The first step was to generate the symbolic constraint in a string presentation. This is done by taking the operands, finding if they are constant or local values, and then converting them to human readable form. The predicate is converted using an enumeration described by the LLVM predicate documentation.

We then insert the function call `@print_string` as seen in line 1 of Code Snippet 3. We can see that `@print_string` takes two arguments `@0` and `@1` which are computed based on the `fcmps` operands. Next we need to determine the value of both operands during the program execution. To do this we pass both operands to logging functions of the appropriate type. As we know the operands are 32 bit floating points we call the `@print_float` function. This function takes in a string and a 32 bit floating point value. We can see that the first call to the logging function in line 2 takes the first operand of the `fcmp` in line 4, while the second call to the logging function prints out the second operand of the `fcmp`. These logging functions would thus print out the both the symbolic constraints as well as the runtime value of each operand.

Code Snippet 3: New Equivalent Instrumented Code

```
1 Inst (call): call void @print_string(i8*...@0, i8*...@1)
2 Inst (call): call void @print_float(i8*...@2, i32...%2)
3 Inst (call): call void @print_float(i8*...@3, i32...0.0)
4 Inst (fcmp): %cmp = fcmp oeq float %2, 0.0
5 Inst (call): call void @print_bool(i8*...@4, i1...%cmp)
6 Inst (br): br i1 %cmp, label %if.then, label %if.else
```

Finally in order to know the result of the compare statement in line 4 of Code Snippet 3, we need to insert one final logging

function. The result of the `fcmp` is a 1 bit integer as seen in line 6. We thus call the appropriate logging function `print_bool` which takes in a 1 bit integer.

6.2 Load Instruction

The first instruction both Algorithm 1 and 2 consider is a load instruction[7]. Each time a load instruction was identified we kept track of which variables were being loaded into new memory locations. A typical example is shown below in Code Snippet 4. In this example you can see that a local variable `%input` is being loaded into another local variable `%2`. The next instruction in line 2, a compare instruction, uses the local variable `%2`. However, in order to print out the constraints for the compare instruction we need to know that the local variable `%2` is in fact `%input`.

Code Snippet 4: Load Instruction

```
1 Inst (load): %2 = load i32, i32* %input, align 4
2 Inst (icmp): %cmp = icmp sgt i32 %2, 0
```

In order to keep track of these mappings we create a vector of pairings. Each pair consists of a memory location and a variable name. For instance in Line 1 when `%input` is stored into `%2`, a mapping between the local variables name `%input` and the memory address of local variable `%2` is created. Thus when the local variable `%2` is used, our algorithm can search through the vector of known mappings and identify the variable name which is being used in any future instruction. Doing this allows us track of variables throughout the program.

6.3 Compare Instruction

Compare instructions, like the one seen line 2 of Code Snippet 4, are points in the program where local variables are compared. When compare instructions are paired with a branch instruction they create branches whose execution depends on the compare instruction result. Our technique thus needs to generate a symbolic constraint for each of the compare instructions. The LLVM pass we implemented checks if the instruction is a compare instruction[5]. We used the base compare instruction which all other compares inherit from, allowing us to keep the LLVM pass general. Our pass then iterates through each of the operands. If the operand is a constant value it is converted to either a constant integer or float depending on the type. If the operand is a local variable if a name can not be generated from the actual operand directly we search through the known local variable mappings. The final step is to determine the type of predicate which is used. We used the enum of predicates defined by LLVM to decode the predicate from an integer into a human readable string.

6.4 Branch Instruction

The last instruction we considered were branch instructions[4]. Branch instructions were only considered by the dynamic pass as we needed to evaluate them at runtime. Code Snippet 5 shows an example of a branch instruction.

Code Snippet 5: Branch Instruction

```
1 Inst (icmp): %2 = icmp sgt i32 %1, 0
2 // Print will be inserted here by logging
3 Inst (br): br i1 %2, label %if.then, label %if.else
```

We can see that we are interested in the value of %2, as this will determine whether the program branches or not. We are able to get the memory location of the local variable %2 by requesting the branch condition and converting it to an LLVM value type. This type is then passed to our logging function which then inserts a print instruction inbetween lines 1 and 3. The logging function will then evaluate the value of %2 at runtime.

7 STUDY

Due to the scope of this project our main goal was to prototype this system to the point where feasibility could be shown. This study aims at showcasing this prototype while showcasing the design feasibility. The dynamic and static pass were crucial to the success of the concolic testing system. Thus both the dynamic and static LLVM pass were tested with the most common numeric data types namely:

- Integers (32 bit and 64 bit)
- Floats (32 bit and 64 bit)
- Booleans (1 bit integers)
- Strings (8 bit integer pointers)

We also tested our LLVM passes on all of the statements which could cause branching such as if statements and loops. Interestingly, loops are handled in the exact same way as if statements at a bitcode level. For example, consider Code Snippet 6. The while loop was converted to an *icmp* statement which is the same instruction type an if statement is converted to. Thus we later learnt that testing on both loops and if statements was actually redundant.

Code Snippet 6: Converting A Loop to Bit Code

```
1 Inst (icmp): %1 = icmp slt i32 %2, %3
2 Inst (br): br i1 %1, label %while.body, label %while.end
```

In order to prototype the system we implemented a constraint solver in Python using Z3. The constraint solver only handled integer values and was implemented in order to allow for a full run of the system which will be discussed later in this section.

We decided that fully automating the approach was not fundamental to showcasing the systems capabilities and thus only automated certain labour parts of the system. For instance, the process of automating the LLVM build, instrumentation, bitcode linking, and running of programs was done using shell scripts. The constraint solver was written in a single Python script and consumed trace files to automate the process of passing constraints to the solver. However in order to run a full test we manually generated traces and passed them to the solver which gave us input we manually passed back to the LLVM.

7.1 Example Execution

In this study, we go through an example using the program in code snippet 7. To better understand how the execution works we also present Figure 2 which shows which paths were executed during the concolic testing.

In our first run, we use the constraints solver to generate a random integer, which in our case was -8361 as seen in line 2 for Code Snippet 8. We then run our program on this input and find that the only constraint we evaluate is *input* > 0 as seen in Code Snippet 8. This input is evaluated as false, and there are no

Code Snippet 7: Original Program

```
1 int main(int argc, char **argv){
2   int input = atoi(argv[1]);
3   int x=0;
4   if(input>0) {
5     if(input > 5) {
6       if(input < 3){
7         x = 2;
8       }
9     }
10    else{
11      x = 1;
12    }
13  }
14  else{
15    x = 3;
16  }
17  return x;
18 }
```

more predicates to evaluate on this path. Code Snippet 8 shows the system output for both the constraint solver and LLVM dynamic pass. We then pass the constraints from the LLVM trace to our constraint solver program, *solveConstraints.py*. This program negates the most recent un-negated constraint. In this case we negate the only constraint namely, *!(input > 0)*. We then check if there is a satisfiable solution, which there is, and then generates new input.

Code Snippet 8: System Output - First Iteration

```
1 How many inputs values does this program take: 1
2 Your values are: [-8361]
3 =====
4 Constraint Hit: input > 0
5 Operand (input): -8361 | Operand (0): 0
6 Predicate Evaluation: False
```

The new input which was generated was the value 1 as seen in line 1 of Code Snippet 9. On the second run, the program returns a trace saying that it evaluated two constraints upon execution, *input* > 0 as true and *input* > 5 as false. Each of these branches can be seen in the second iteration Figure 2 as well as in the LLVM dynamic pass output in Code Snippet 9. When we pass these constraints to our input generator, it negates *input* > 5 and generates the input of 6.

Code Snippet 9: System Output - Second Iteration

```
1 sat - New Input: [a = 1]
2 =====
3 Constraint Hit: input > 0
4 Operand (input): 1 | Operand (0): 0
5 Predicate Evaluation: True
6 =====
7 Constraint Hit: input > 5
8 Operand (input): 1 | Operand (5): 5
9 Predicate Evaluation: False
```

The third run, as seen in Code Snippet 10, evaluates three constraints, *input* > 0 as true, *input* > 5 as true and *input* < 3 as false. We pass these constraints to our constraint solver with the last constraint, *!(input < 3)*, negated. Since input can not be both greater than 5 and less than 3, this returns unsat. It is at this point that we have executed all paths in the control flow graph.

This ability to both stop after an unsat condition was found and consider flow sensitive symbolic constraints is extremely useful. Consider for instance our static analysis LLVM pass whose results for the same original program are shown in Code Snippet 11. There

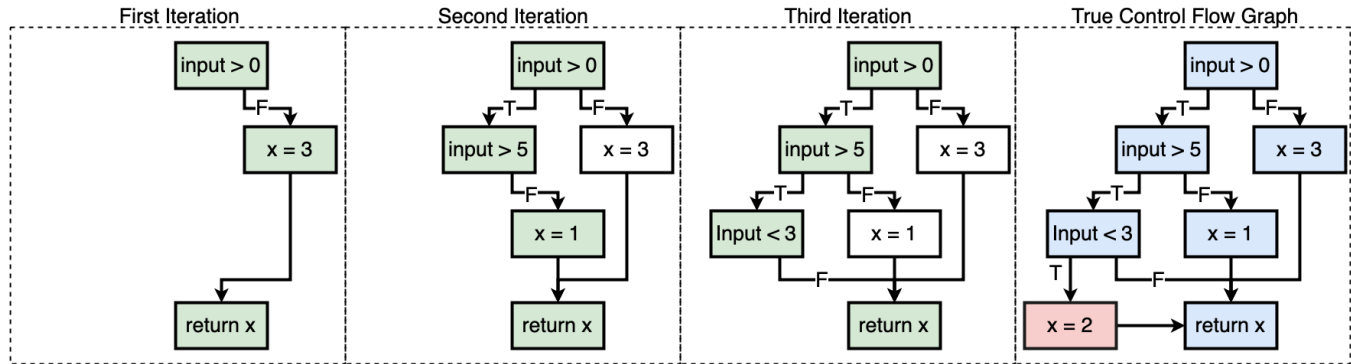


Figure 2: The executed paths (shown in green) for each iteration of our study. The true control flow graph is depicted on the right with blue nodes depicting both nodes which our execution engine found and red nodes showing nodes which our execution engine found to be unsatisfiable.

is no way to tell which constraints are along the same branch, and the order in which these constraints are applied are unknown.

Code Snippet 10: System Output - Final Iteration

```

1 sat - New Input: [a = 6]
2 =====
3 Constraint Hit: input > 0
4 Operand (input): 6 | Operand (0): 0
5 Predicate Evaluation: True
6 =====
7 Constraint Hit: input > 5
8 Operand (input): 6 | Operand (5): 5
9 Predicate Evaluation: True
10 =====
11 Constraint Hit: input < 3
12 Operand (input): 6 | Operand (3): 3
13 Predicate Evaluation: False

```

Code Snippet 11: System Output - Static Analysis

```

1 Constraint: input > 0
2 Constraint: input > 5
3 Constraint: input < 3

```

Similarly our engines ability to determine that the constraints are unsat is extremely useful. Consider the case where path coverage is being used as a stopping condition for a random input fuzzer. The fuzzer would waste tons of execution trying to hit a branch which after only three executions of this approach was determined to be unsat.

8 CONCLUSION

In this paper, we first presented and implemented a framework for doing concolic testing on LLVM bitcode. This involves two LLVM passes and a python program. The first LLVM pass generates symbolic constraints for a program statically. The second LLVM pass inserts prints statements that generates program traces of the constraints evaluated on that execution. These constraints are then passed to the python program which uses an SMT solver to find new inputs that satisfies new constraints of a program.

When compared against the constraints generated by the static pass, the dynamic pass along with constraint solver generate fewer inputs that still satisfy all paths. This is due to the fact that the dynamic pass is flow sensitive unlike the static pass. If we were to generate input based solely on the static constraints, we would

have to generate an input for every satisfiable conjunction of the constraints to ensure we hit all paths.

REFERENCES

- [1] Anders Moller. 2019. TIP. <https://github.com/cs-au-dk/TIP>. [Online; accessed 7-December-2019].
- [2] Cristian Cadar, Daniel Dunbar, Dawson R Engler, et al. 2008. KLEE: Unassisted and Automatic Generation of High-Coverage Tests for Complex Systems Programs.. In *OSDI*, Vol. 8. 209–224.
- [3] Leonardo de Moura and Nikolaj Bjørner. 2008. Z3: An Efficient SMT Solver. In *Tools and Algorithms for the Construction and Analysis of Systems*, C. R. Ramakrishnan and Jakob Rehof (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 337–340.
- [4] LLVM Project. 2019. BranchInst Class Reference. http://llvm.org/doxygen/classllvm_1_1BranchInst.html. [Online; accessed 8-December-2019].
- [5] LLVM Project. 2019. CmpInst Class Reference. http://llvm.org/doxygen/classllvm_1_1CmpInst.html. [Online; accessed 8-December-2019].
- [6] LLVM Project. 2019. LLVM's Analysis and Transform Passes - LLVM 10 documentation. <https://github.com/cs-au-dk/TIP>. [Online; accessed 8-December-2019].
- [7] LLVM Project. 2019. LoadInst Class Reference. http://llvm.org/doxygen/classllvm_1_1LoadInst.html. [Online; accessed 8-December-2019].
- [8] Matthew Dwyer. 2019. TIP. <https://github.com/matthewbdwyer/tip/>. [Online; accessed 7-December-2019].
- [9] Yannic Noller, Rodry Kersten, and Corina S. Păsăreanu. 2018. Badger: Complexity Analysis with Fuzzing and Symbolic Execution. In *Proceedings of the 27th ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA 2018)*. ACM, New York, NY, USA, 322–332. <https://doi.org/10.1145/3213846.3213868>
- [10] Nick Stephens, John Grosen, Christopher Salls, Andrew Dutcher, Ruoyu Wang, Jacopo Corbetta, Yan Shoshitaishvili, Christopher Kruegel, and Giovanni Vigna. 2016. Driller: Augmenting Fuzzing Through Selective Symbolic Execution.. In *NDSS*, Vol. 16. 1–16.
- [11] Xinyu Wang, Jun Sun, Zhenbang Chen, Peixin Zhang, Jingyi Wang, and Yun Lin. 2018. Towards Optimal Concolic Testing. In *Proceedings of the 40th International Conference on Software Engineering (ICSE '18)*. ACM, New York, NY, USA, 291–302. <https://doi.org/10.1145/3180155.3180177>